

# CS553 Homework #5

---

## Sort on Hadoop/Spark

### Instructions:

- *Maximum Points: 100%*
- *Please post your questions to the Piazza forum*
- *Only a softcopy submission is required; it will automatically be collected through GIT after the deadline; email confirmation will be sent to your HAWK email address*
- *Late submission will be penalized at 10% per day; an email to the TA with the subject “CS553: late homework submission” must be sent*

### 1. Introduction

The goal of this programming assignment is to enable you to gain experience programming with:

- The Hadoop framework (<http://hadoop.apache.org/>)
- The Spark framework (<http://spark.apache.org/>)

In Homework #6, you implemented an external sort and compared it to the Linux sort. You will now expand into implementing sort with Hadoop and with Spark.

### 2. Your Assignment

This programming assignment covers sort through Hadoop and Spark on multiple nodes. You must use a Chameleon node using Bare Metal Provisioning (<https://www.chameleoncloud.org>). You must deploy Ubuntu Linux 16.04 using either “compute-skylake” or “compute-haswell” nodes, at either UC or TACC sites. Once you create a lease (up to 7 days are allowed), and start your 1 physical node, and Linux boots, you will find yourself with a physical node with 24 CPU cores, 48 hardware threads, 128GB to 192GB of memory (depending on if its Haswell or Skylake), and 250GB SSD hard drive. You will install your favorite virtualization or containerization tools (e.g. virtualbox, KVM, qemu, Docker, lxc/lxd), and use it to deploy two different VMs/containers with the following sizes: small.instance (4-cores, 8gb ram, 50gb disk), and large.instance (16-cores, 32GB ram, 200gb disk).

This assignment will be broken down into several parts, as outlined below:

**Hadoop File System and Hadoop Install:** Download, install, configure, and start the HDFS system (that is part of Hadoop, <https://hadoop.apache.org>) on a virtual cluster with 1 large.instance, and then again on a virtual cluster with 4 small.instances. You must turn off replication, or you won’t have enough storage capacity to conduct your experiments.

**Datasets:** Once HDFS is operational, you must generate your dataset with gensort (<http://www.ordinal.com/gensort.html>); you will create 4 workloads: data-1GB, data-4GB, data-16GB, and data-64GB. You may not have enough room to store them all, and run your compute workloads. Make sure to cleanup after each run. Remember that you will typically need 3X the storage, as you have the original input data, temporary data, and output data. Configure Hadoop to run on the virtual cluster, on 1 large.instance and 4 small.instances. You may need a 5<sup>th</sup> virtual machine to run parts of Hadoop (e.g. name node, scheduler, etc).

**Spark Install:** Download, install, configure, and start Spark (<https://spark.apache.org>). Note that you will need the HDFS installation for Hadoop to work from and to.

**Shared Memory Sort:** Run your shared memory sort from HW5 on the small.instance and large.instance defined above on the different datasets.

**Linux Sort:** Run the Linux Sort on the small.instance and large.instance defined above on the different datasets.

**Hadoop Sort:** Implement the HadoopSort application (you must use Java). You must specify the number of reducers to ensure you use all the resources of the 4-node cluster. You must read the data from HDFS, sort it, and then store the sorted data in HDFS and validate it with valsort. Measure the time from reading from HDFS, sorting, and writing to HDFS; do not include the time to run valsort.

**Spark Sort:** Implement the SparkSort application (you must use Java). Make sure to use RDD to speed up the sort through in-memory computing. You must read the data from HDFS, make use of RDD, sort, and write the sorted data back to HDFS, and finally validate the sorted data with valsort. Measure the time from reading from HDFS, sorting, and writing to HDFS; do not include the time to run valsort.

**Performance:** Compare the performance of your shared-memory external sort, the Linux “sort” (more information at <http://man7.org/linux/man-pages/man1/sort.1.html>), Hadoop Sort, and Spark Sort on a 4-node cluster with the 1GB, 4GB, 16GB, and 64GB datasets. Fill in the table below, and then derive new tables or figures (if needed) to explain the results. Your time should be reported in milliseconds.

Complete Table 1 outlined below. Perform the experiments outlined above, and complete the following table:

*Table 1: Performance evaluation of sort (report time to sort in milliseconds)*

Experiment	Shared Memory	Linux	Hadoop Sort	Spark Sort
1 small.instance, 1GB dataset				
1 small.instance, 4GB dataset				
1 small.instance, 16GB dataset				
1 large.instance, 1GB dataset				
1 large.instance, 4GB dataset				
1 large.instance, 16GB dataset				
1 large.instance, 64GB dataset				
4 small.instances, 1GB dataset	N/A	N/A		
4 small.instances, 4GB dataset	N/A	N/A		
4 small.instances, 16GB dataset	N/A	N/A		
4 small.instances, 64GB dataset	N/A	N/A		

Some of the things that will be interesting to explain are: how many threads, mappers, reducers, you used in each experiment; how many times did you have to read and write the dataset for each experiment; what speedup and efficiency did you achieve?

For the 64GB workload (with both 1 large instance and 4 small instances), monitor the disk I/O speed (in MB/sec), memory utilization (GB), and processor utilization (%) as a function of time, and generate a plot for the entire experiment. Here is an example of a plot that has cpu utilization and memory utilization (<https://i.stack.imgur.com/dmYAB.png>), plot a similar looking graph but with the disk I/O data as well as a 3<sup>rd</sup> line. Do this for both shared memory benchmark (your code) and for the Linux Sort. You might find some online info useful on how to monitor this type of information (<https://unix.stackexchange.com/questions/554/how-to->

[monitor-cpu-memory-usage-of-a-single-process](#)). For multiple instances, you will need to combine your monitor data to get an aggregate view of resource usage. Do this for all four versions of your sort. After you have all eight graphs (2 system configurations and 4 different sort techniques), discuss the differences you see, which might explain the difference in performance you get between the two implementations. Make sure your data is not cached in the OS memory before you run your experiments.

Note that you do not have to artificially limit the amount of memory your sort can use as the VMs will be configured with a limited amount of memory. What conclusions can you draw? Which seems to be best at 1 node scale (1 large.instance)? Is there a difference between 1 small.instance and 1 large.instance? How about 4 nodes (4 small.instance)? What speedup do you achieve with strong scaling between 1 to 4 nodes? What speedup do you achieve with weak scaling between 1 to 4 nodes? How many small.instances do you need with Hadoop to achieve the same level of performance as your shared memory sort? How about how many small.instances do you need with Spark to achieve the same level of performance as you did with your shared memory sort? Can you draw any conclusions on the performance of the bare-metal instance performance from HW5 compared to the performance of your sort on a large instance through virtualization?

Can you predict which would be best if you had 100 small.instances? How about 1000? Compare your results with those from the Sort Benchmark (<http://sortbenchmark.org>), specifically the winners in 2013 and 2014 who used Hadoop and Spark. Also, what can you learn from the CloudSort benchmark, a report can be found at ([http://sortbenchmark.org/2014\\_06\\_CloudSort\\_v\\_0\\_4.pdf](http://sortbenchmark.org/2014_06_CloudSort_v_0_4.pdf)).

### 3. What you will submit

The grading will be done according to the rubric below:

- Hadoop (installation/config) sort implementation/scripts: 20 points
- Spark (installation/config) sort implementation/scripts: 20 points
- Performance evaluation (data): 30 points
- Performance evaluation (Q&A and explanations): 30 points

The maximum score that will be allowed is 100 points.

You are to write a report (hw6\_report.pdf). Add a brief description of the problem, methodology, and runtime environment settings. You are to fill in the table on the previous page. Please explain your results, and explain the difference in performance? Include logs from your application as well as valsort (e.g. standard output) that clearly shows the completion of the sort invocations with clear timing information and experiment details; include separate logs for shared memory sort, Linux sort, Hadoop sort, and Spark sort, for each dataset. Valsort can be found as part of the gensort suite (<http://www.ordinal.com/gensort.html>), and it is used to validate the sort. As part of your submission you need to upload to your private git repository your run scripts, build scripts, the source code for your implementation (shared memory sort, Hadoop sort, and spark sort), the report, a readme file (with how to build and use your code), and several log files. Make sure to answer all the questions posed from Section 2. Some of your answers might require a graph or table with data to substantiate your answer.

Here are the naming conventions for the required files:

- |  |                      |
|--|----------------------|
| • Makefile / build.xml (Ant) / pom.xml (Maven) | • readme.txt         |
| • MySort.java / mysort.c / MySort.cpp          | • mysort64GB.log     |
| • HadoopSort.java                              | • linsort64GB.log    |
| • SparkSort.java                               | • hadoopsort64GB.log |
| • Scripts                                      | • sparksort64GB.log  |
| • Hw6_report.pdf                               |                      |

When you have finished implementing the complete assignment as described above, you should submit your solution to your private git repository. Each program must work correctly and be detailed in-line documented.

**Submit code/report through GIT.** If you cannot access your repository contact the TAs. You can find a git cheat sheet here: <https://www.git-tower.com/blog/git-cheat-sheet/>

**Grades for late programs will be lowered 10% per day late.**